

CSE231 – Operating Systems

Assignment-3

Name: Ananya Lohani
Roll Number: 2019018
Branch: CSE

Description:

In this assignment, I have modified the `sched_entity` structure to contain a new field `rt_nice` which is the soft real-time requirement of a process (a value of `x` as `rt_nice` means that the process must receive at least `x` units of time-slice). To give priority to processes with soft real-time requirements I have modified the CFS scheduler to schedule tasks on the basis of `rt_nice` first, and `vruntime` second.

rt_nice System Call implementation:

I added a new field `rt_nice` to the `sched_entity` struct in `/kernel/sched/sched.h` and initialised it to 0 in the `__sched_fork()` function in `/kernel/sched/core.c`. I have added a system call `rt_nice` which takes two arguments: **PID of the process** and the **soft real-time requirement**. The system call gets the task corresponding to the given PID and sets the `rt_nice` field of its `sched_entity` to the given value.

Modified CFS implementation:

The CFS scheduler of Linux has been modified to give priority to the `rt_nice` values before giving priority to the `vruntime` of processes in the run-queue.

Functions modified:-

- **entity_before():** This function acts as a comparator between the `vruntime` of 2 `sched_entity` struct pointers. I added the code that compares the `rt_nice` values of the pointers first, and if both pointers' `rt_nice` is equal to 0, it compares their `vruntime`.
- **__pick_next_entity():** This function picks the next process to be executed. I modified it to check if any process in the runqueue has a non-zero `rt_nice` value and selected the process with the minimum

`rt_nice` to be executed. If no process has `rt_nice` greater than zero, it will execute the next process according to `vruntime`.

- **update_curr():** If the `rt_nice` value of a `sched_entity` is greater than 0, it updates `rt_nice` by subtracting from it the amount of time that the process ran.

Errors handled:-

Errors returned by the system call are handled by `perror()` in test files.

- **EINVAL:** If an invalid pid was given(pid not between 1 and 2147483647) or if invalid value for `rt_nice` was given(<0).
- **ESRCH:** If there's no process corresponding to the given pid.

Testing the Scheduler:

The working of the scheduler can be tested by the test files provided: `test1.c` and `test2.c`. `test1.c` takes a command line argument `rt_nice` from the user and calls the `rt_nice` system call to assign that value to itself. It runs a loop of 2000000000 and calculates internally the time it took to finish the process. `test2.c` forks a child process and assigns `rt_nice` as 200 to the parent and 0 to the child. Both the parent and the child processes execute a loop of 300000000 and calculate the time it took to finish the loop.

Expected output:

In `test1.c`, the process with a positive value of `rt_nice` took less time to be completed. The expected output is as follows:

```
→ modified-cfs-scheduler git:(main) x ./test1
rt_nice: 0 sum = 2000000000 time: 6.015671 seconds
→ modified-cfs-scheduler git:(main) x ./test1 100
rt_nice: 100 sum = 2000000000 time: 5.950315 seconds
```

In `test2.c`, the parent process with `rt_nice` = 200 finished the loop faster than the child process with `rt_nice` = 0. The expected output is as follows:

```
./test2
Parent Process PID: 4467, rt_nice = 200, time: 0.085729 seconds
Child Process PID: 4468, rt_nice = 0, time: 0.090100 seconds
```